



RUHR-UNIVERSITÄT BOCHUM

Methods of User Authentication: Programming Tutorial 1

Maximilian Golla

Do We Care?



Schneier on Security

Essays

News



← Information in Your Boarding Pass's Bar Code

I'm a Guest on "Adam Ruins Everything" →

Crypto

Speaking

SHA-1 Freestart Collision

Newsletter

Blog

There's a new <u>cryptanalysis result</u> against the hash function SHA-1:

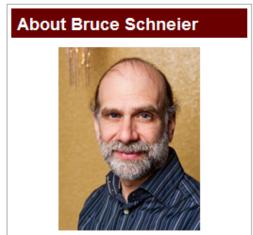
Books

Abstract: We present in this article a freestart collision example for SHA-1, i.e., a collision for its internal compression function. This is the first practical break of the full SHA-1, reaching all 80 out of 80 steps, while only 10 days of computation on a 64 GPU cluster were necessary to perform the attack. This work builds on a continuous series of cryptanalytic advancements on SHA-1 since the theoretical collision attack breakthrough in 2005. In particular, we extend the recent freestart collision work on reduced-round SHA-1 from CRYPTO 2015 that leverages the computational power of graphic cards and adapt it to allow the use of boomerang speed-up techniques. We also leverage the cryptanalytic techniques by Stevens from EUROCRYPT 2013 to obtain optimal attack conditions, which required further refinements for this work. Freestart collisions, like the one presented here, do not directly imply a collision for SHA-1.

However, this work is an important milestone towards an actual SHA-1 collision and it further shows how graphics cards can be used very efficiently for these kind of attacks. Based on the state-of-the-art collision attack on SHA-1 by Stevens from EUROCRYPT 2013, we are able to present new projections on the computational/financial cost required by a SHA-1 collision computation. These projections are significantly lower than previously







Recap: Cryptographic Hash Functions



| German | English | Given | Should be Hard |
|---------------------------------|---|---------|-------------------------------------|
| Einwegfunktion | Preimage resistance | hash(X) | X |
| Schwache Kollisionsresistenz | 2 nd -Preimage resistance | Y | Y' with hash(Y) == hash(Y') |
| Starke Kollisionsresistenz | Collision resistance | - | Z,Z' with hash(Z) == hash(Z') |

Recap: The Problem



Hash functions are "fast" and deterministic

\$> echo "mySecurePassword" -n | sha1sum 07b5e42f1eebdafc0591f67579a9db194e2c8f97

\$> echo "mySecurePassword1" -n | sha1sum 53aeab619c8fefd80c01b81451b9c1d91ce09f7f

\$> echo "mySecurePassword" -n | sha1sum 07b5e42f1eebdafc0591f67579a9db194e2c8f97

Recap: Storing Passwords



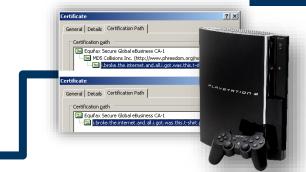
- We do not store plaintext passwords but,
 - salted (e.g., 128-bit)
 - **iterated** (e.g., $2^{12} = 4096$ rounds)
 - hashed (e.g., bcrypt / scrypt / PBKDF2 / Argon2i)
 passwords in the database.
- Output string containing everything we need to know:

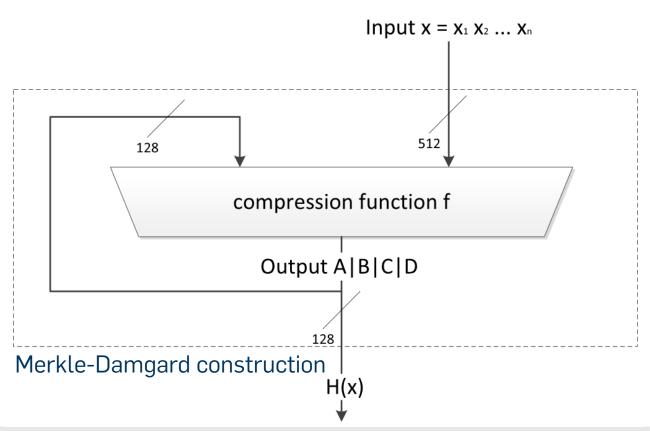
\$2a\$12\$8vxYfAWCXe0Hm4gNX8nzwuqWNuk0kcMJ1a9G2tD71ipotEZ9f80Vu

Example MD5

5Sahre
RUB

- Ronald L. Rivest 1992, successor of MD4
- Collision resistance 2⁶⁴ (Birthday Attack)
- Collision in 2004 (Wang et al.)
- Highly practical in 2008 (Stevens et al.)





In More Detail



- 64 rounds (4x 16)
- Input: 512-bit, Output: 128-bit

Every single round consist of:


- √ 4x 32-bit variables

 A

 B

 C

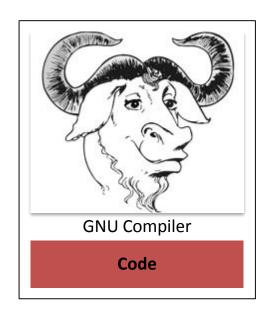
 D
- ✓ M_i Message (Input string)
- \checkmark K_i Constants (By the author)
- Non-linear compression function Boolean bitwise functions (AND, OR, XOR, NOT)
- ✓ \bigcirc Modular addition (2³²)





A single round of MD5



Optimizing Hash-Algorithms as an Attacker







How to Speed Up!



- 1. Message padding
 - Input x (Password candidate) has length b = |x|
 - Append single bit "1"

-> Use fixed padding

- Fill with "0"-bits until 448 mod 512
- Add length b as two 32-bit words (lower order first)
- 2. Defining constants

a0 := 0x67452301 For all i from 0 to 63:

CO := 0x98BADCFE

d0 := 0x10325476 -> Precompute

3. Compression functions

 $F(X,Y,Z) = (X \land Y) \lor (\neg X \land Z)$ $G(X,Y,Z) = (X \land Z) \lor (Y \land \neg Z)$ $H(X,Y,Z) = X \oplus Y \oplus Z$

 $I(X, I, Z) = X \oplus I \oplus Z$

 $I(X, Y, Z) = Y \oplus (X \vee \neg Z)$

-> Use optimized functions

($0 \le i \le 15$): F := D xor (B and (C xor D))

 $(16 \le i \le 31)$: F := C xor (D and (B xor C))







Constants

Is a variable whose value cannot be changed once it is initially bound to a value. There's usually a small performance and memory advantage in using constants.

```
const __m128i K00_19 = SET1INT(0x5A827999);
```

Multiple Variable Declaration

All variables are allocated in one memory space. The count of variables can impact the performance.

```
uint32_t interm, interm2, interm3, interm4, foo, bar;
```



Static Inline Functions

Making a function an inline function suggests that calls to the function be as fast as possible. The exact situations under which an inline function actually gets in-lined vary depending on the compiler; the details of this are complicated.

```
static inline __m128i F00_19(__m128i B, __m128i C, __m128i D) {
   return XOR(D, AND(B, XOR(C, D)));
}
```

Macros

They are a way of eliminating function call overhead because they are always expanded in-line, unlike the "inline" keyword which is an often-ignored hint to the compiler, and didn't even exist (in the standard) prior to C99.

```
#define ROTATE_LEFT(x, n) (((x) << (n)) | ((x) >> (32-(n))))
```



Inline Assembly

With inline assembly we are able to further accelerate the execution of applications by handcrafting the assembler codes for the most performance-sensitive parts.

```
__asm__("ROR %1,30\n\t" : "=a" (var) : "a" (var));
```

FIPS PUB 180-1 Alternative F-Functions

Equivalent expressions that reduce the necessary amount of instructions per invocation.

```
(0 ≤ i ≤ 19): // Original
    f = (b and c) or ((not b) and d)
(0 ≤ i ≤ 19): // Alternative
    f = d xor (b and (c xor d))
```



Loop unrolling

```
int i = 0;
for (; i < 16; i++) {
    x[i] = buffer[i * 4 + 3] << 24 | buffer[i * 4 + 2] << 16 | buffer[i * 4 + 1] << 8 | buffer[i * 4 + 0];
/* Unrolled */
x[0] = buffer[0 * 4 + 3] << 24
                                  buffer[ 0 * 4 + 2] << 16
                                                             buffer[ 0 * 4 + 1] << 8
                                                                                       buffer[ 0 * 4 + 0];
x[1] = buffer[1 * 4 + 3] << 24
                                  buffer[ 1 * 4 + 2] << 16
                                                             buffer[ 1 * 4 + 1] << 8
                                                                                       buffer[ 1 * 4 + 0];
x[2] = buffer[2 * 4 + 3] << 24
                                  buffer[ 2 * 4 + 2] << 16
                                                             buffer[ 2 * 4 + 1] << 8
                                                                                       buffer[ 2 * 4 + 0];
                                                                                       buffer[ 3 * 4 + 0];
x[3] = buffer[3 * 4 + 3] << 24
                                  buffer[ 3 * 4 + 2] << 16
                                                             buffer[ 3 * 4 + 1] << 8
x[4] = buffer[4 * 4 + 3] << 24
                                  buffer[ 4 * 4 + 2] << 16
                                                             buffer[ 4 * 4 + 1] << 8
                                                                                       buffer[ 4 * 4 + 0];
x[5] = buffer[5 * 4 + 3] << 24
                                                             buffer[ 5 * 4 + 1] << 8
                                  buffer[ 5 * 4 + 2] << 16
                                                                                       buffer[5*4+0];
x[6] = buffer[6 * 4 + 3] << 24
                                  buffer[ 6 * 4 + 2] << 16
                                                             buffer[ 6 * 4 + 1] << 8
                                                                                       buffer[6*4+0];
                                  buffer[ 7 * 4 + 2] << 16
x[7] = buffer[7 * 4 + 3] << 24
                                                             buffer[ 7 * 4 + 1] << 8
                                                                                       buffer[ 7 * 4 + 0];
                                  buffer[ 8 * 4 + 2] << 16
x[8] = buffer[8 * 4 + 3] << 24
                                                             buffer[ 8 * 4 + 1] << 8
                                                                                       buffer[8 * 4 + 0];
x[9] = buffer[9 * 4 + 3] << 24
                                  buffer[ 9 * 4 + 2] << 16
                                                             buffer[ 9 * 4 + 1] << 8
                                                                                       buffer[ 9 * 4 + 0];
x[10] = buffer[10 * 4 + 3] << 24
                                  buffer[10 * 4 + 2] << 16
                                                             buffer[10 * 4 + 1] << 8
                                                                                       buffer[10 * 4 + 0];
x[11] = buffer[11 * 4 + 3] << 24
                                  buffer[11 * 4 + 2] << 16
                                                             buffer[11 * 4 + 1] << 8
                                                                                       buffer[11 * 4 + 0];
x[12] = buffer[12 * 4 + 3] << 24
                                  buffer[12 * 4 + 2] << 16
                                                             buffer[12 * 4 + 1] << 8
                                                                                       buffer[12 * 4 + 0];
x[13] = buffer[13 * 4 + 3] << 24
                                  buffer[13 * 4 + 2] << 16
                                                             buffer[13 * 4 + 1] << 8
                                                                                       buffer[13 * 4 + 0];
x[14] = buffer[14 * 4 + 3] << 24
                                  buffer[14 * 4 + 2] << 16
                                                             buffer[14 * 4 + 1] << 8
                                                                                       buffer[14 * 4 + 0];
x[15] = buffer[15 * 4 + 3] << 24
                                  buffer[15 * 4 + 2] << 16
                                                             buffer[15 * 4 + 1] << 8
                                                                                       buffer[15 * 4 + 0];
```



#pragma GCC optimize("unroll-loops")

Framework



Can be found online in Moodle

```
shal.h, testbench.c, testbench.h
(Do not touch, do not submit)
```

```
SHA-1 Team-01 1.c
```

- Teams of two students
- Rename file according to scheme
- Write your complete code in this file

```
#include "shal.h"

int crackHash(struct state hash, char *result) {
    result[0] = 'a';
    result[1] = 'b';
    result[2] = 'c';
    result[3] = 'd';
    result[4] = 'e';
    result[5] = 'f';
    /* Found */
    return(EXIT_SUCCESS);
    /* Not found */
    return(EXIT_FAILURE);
}
```

```
SHA1 Coding Framework

Type

286 bytes C header

SHA-1_Team-XX_1.c 770 bytes C source code

testbench.c 1,5 kB C source code

testbench.h 478 bytes C header
```

```
struct state {
    uint32_t a;
    uint32_t b;
    uint32_t c;
    uint32_t d;
    uint32_t e;
};
```

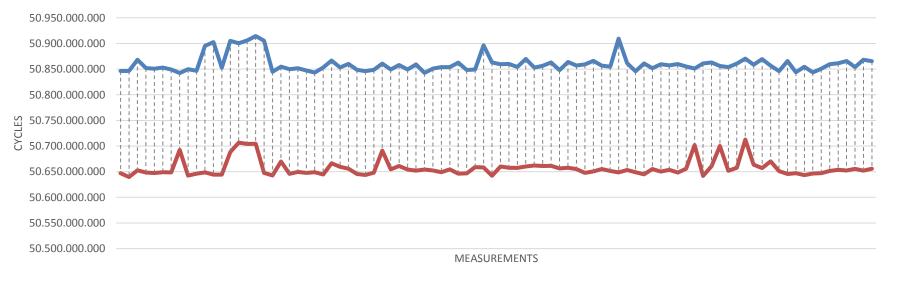




Performance measurement via execution-time and CPU-cycles

```
clock_t start = clock();
uint64_t cnt = rdtsc();
error = crackHash(hash, result);
cnt = rdtsc() - cnt;
clock_t stop = clock();
double elapsed = (double)(stop - start) * 1000.0 / CLOCKS_PER_SEC;
```





Framework



Example compile and execute

```
$> gcc -02 -Wall -fomit-frame-pointer -msse2 -masm=intel
testbench.c <YOUR IMPL>.c -o crackSHA1
```

\$> ./crackSHA1 68d8572c2662b0f06f723d7d507954fb038b8558

Hash: 0x68D8572C 2662B0F0 6F723D7D 507954FB 038B8558

Preimge: aaabbb

Cycles: 39452294775

Time: 14633.591000

Sneak Peek



- Next week
 - Early-Exit
 - Reverse Computation
 - Message Expansion Attack by Jens Steube (SHA-1 specific)
 - **–** ..
 - SIMD via Intel Streaming SIMD Extensions 2 (SSE2)

Happy Coding ©

First Round of Bonus Points - Submission Deadline: 13. November 2015 - 11:59 p.m. Error and Warning Free!!!